



Toronto MongoDB Atlas Workshop

Eugene Kang, Sr. Solutions Architect, MongoDB
Jeegar Ghodasara, Sr. Solutions Architect, MongoDB
Will Chow, Sr. Solutions Architect, MongoDB
Amy Rosenberg, Manager, Field Marketing, MongoDB



Agenda




- Introductions and Kickoff
- MongoDB Technical Overview
- Atlas Overview and Demo
- CRUD Operations and Indexes/Explain Plans
- Aggregation Framework
- MongoDB Stitch
- MongoDB Charts
- Closing Remarks

MongoDB Technical Overview



Why MongoDB?

Intelligent Operational Data Platform



**Best way to work
with data**



**Intelligently put data
where you need it**



**Freedom
to run anywhere**



Cases for Change

MongoDB is a modern, operational database that supports a polyglot data strategy – on-premise and in the cloud. This allows us to drive several business critical topics with our customers.

Cloud Data Strategy

Leveraging the right data platforms as part of your overall cloud strategy helps to avoid vendor lock-in.

Legacy Modernization

Current legacy technology stacks can't cope with the range of new business requirements – we can help you modernise in a highly efficient and effective way

Mainframe Offloading

Reduce cost and MIPS on legacy mainframes and enable data to be leveraged for new use cases.

Single View

Provide a holistic view of data entities (e.g. customer) across multiple underlying, disconnected source systems.

Operational Intelligence

Solving the problem of deriving value from existing EDW or Hadoop-based data lake solutions in real-time

Internet of Things (IoT)

MongoDB can help you overcome Scalability & Performance issues that are not being met by many current IoT solutions.

Distributed Ledger/Blockchain

MongoDB is the ideal database/persistence layer for enterprise-grade Distributed Ledger applications


Leading modern general purpose database

 Creative Cloud	 HR Mobile Application	 Real-Time Travel Search	 Mobile Drug Applications	 Mobile Banking	 Multi-Screen TV
 Internet of Things Platform	 Predictive Messaging	 Background Checks as a Service	 Order Capture	 Cryptocurrency Platform	 Logistics Modernization
 E-Commerce System	 Entry Decision System	 Social Security Benefits Program	 Mobile App for Patient Data	 Product Catalog	 Gaming Platform
 Single View of Patient	 Genetic Analysis	 Online Lending	 Online Tax Returns	 Swap Equities Management	 Real-Time Analytics
 E-Commerce Personalization	 E-Commerce Platform	 Marketing Cloud	 Video Streaming	 Log Metadata Store	 Social Media Management
 Website Platform	 Product Catalog	 Identity Theft Protection	 Mass Spectrometer Instrumentation	 Subscriber QoS	 Shopping Cart



Some Terminology


A comparison



** Proper document schema design yields more entity data per document than found in a relational database row



Easy: Contrasting data models



Tabular (Relational) Data Model

Related data split across multiple records and tables

```
{
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),
  "name" : {
    "first" : "John",
    "last" : "Doe" },
  "address" : [
    {
      "location" : "work",
      "address" : {
        "street" : "16 Hatfields",
        "city" : "London",
        "postal_code" : "SE1 8DJ" },
      "geo" : { "type" : "Point", "coord" : [
        51.5065752,-0.109081] } },
    + { ... }
  ],
  "phone" : [
    {
      "location" : "work",
      "number" : "+44-1234567890" },
    + { ... }
  ],
  "dob" : ISODate("1977-04-01T05:00:00Z"),
  "retirement_fund" : NumberDecimal("1292815.75")
}
```

Document Data Model

Related data contained in a single, rich document



Flexible: Adapt to change


```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

```
{  
  "_id" : ObjectId("5ad88534e3632e1a35a58d00"),  
  "name" : {  
    "first" : "John",  
    "last" : "Doe" },  
  "address" : [  
    { "location" : "work",  
      "address" : {  
        "street" : "16 Hatfields",  
        "city" : "London",  
        "postal_code" : "SE1 8DJ"},  
        "geo" : { "type" : "Point", "coord" : [  
          51.5065752,-0.109081]}},  
    + {...}  
  ],  
  "phone" : [  
    { "location" : "work",  
      "number" : "+44-1234567890"},  
    + {...}  
  ],  
  "dob" : ISODate("1977-04-01T05:00:00Z"),  
  "retirement_fund" : NumberDecimal("1292815.75")  
}
```

Add new fields dynamically at runtime



Flexible: Govern



JSON Schema

Enforces strict schema structure over a complete collection for data governance & quality

- Builds on document validation introduced by restricting new content that can be added to a document
- Enforces presence, type, and values for document content, including nested array
- Simplifies application logic

Tunable: enforce document structure, log warnings, or allow complete schema flexibility


Queryable: identify all existing documents that do not comply

Schema Validation Example


```
db.runCommand({  
    collMod: "orders",  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            additionalProperties: false,  
            required: ["item", "price"],  
            properties: {  
                _id: {},  
                item: {  
                    bsonType: "string",  
                    description: "'item' must be a string and is required"  
                },  
                price: {  
                    bsonType: "decimal",  
                    description: "'price' must be a decimal and is required"  
                },  
                quantity: {  
                    bsonType: ["int", "long"],  
                    minimum: 1,  
                    maximum: 100,  
                    exclusiveMaximum: true,  
                    description:  
                        "'quantity' must be short or long integer between 1 and 99"  
                }  
            }  
        }  
    }  
});
```




Versatile: Multiple data models, rich query functionality



JSON Documents



Tabular




Key-Value



Text



Geospatial



Graph

Rich Queries

Point | Range | Geospatial | Faceted Search | Aggregations | JOINS | Graph Traversals



Fully Indexable


Fully featured secondary indexes

Index Types

- Primary Index
 - Every Collection has a primary key index
- Compound Index
 - Index against multiple keys in the document
- MultiKey Index
 - Index into arrays
- Text Indexes
 - Support for text searches
- GeoSpatial Indexes
 - 2d & 2dSphere indexes for spatial geometries
- Hashed Indexes
 - Hashed based values for sharding

Index Features

- TTL Indexes
 - Single Field indexes, when expired delete the document
- Unique Indexes
 - Ensures value is not duplicated
- Partial Indexes
 - Expression based indexes, allowing indexes on subsets of data
- Case Insensitive Indexes
 - supports text search using case insensitive search
- Sparse Indexes
 - Only index documents which have the given field



Easy: MongoDB Multi-Document ACID Transactions



Just like relational transactions

- Multi-statement, familiar relational syntax
- Easy to add to any application
- Multiple documents in 1 or many collections and databases

ACID guarantees

- Snapshot isolation, all or nothing execution
- No performance impact for non-transactional operations

Schedule


- MongoDB 4.0: replica set
- MongoDB 4.2: extended to sharded clusters



Aggregations

Advanced data processing pipeline for transformations and analytics

- Multiple stages
- Similar to a unix pipe
 - Build complex pipeline by chaining commands together
- Rich Expressions





Aggregation Features

A feature rich analytical framework


Pipeline Stages

- \$match
- \$group
- \$facet
- \$geoNear
- \$graphLookup
- \$lookup
- \$project
- \$sort
- \$unwind

Operators

- Mathematical
 - \$add, \$abs, \$subtract, \$multiply, \$divide, \$log, \$log10, \$stdDevPop, \$stdDevSam, \$avg, \$sqrt, \$pow, \$sum, \$zip, \$convert, etc.
- Array
 - \$push, \$reduce, \$reverseArray, \$addToSet, \$arrayElemAt, \$slice, etc.
- Mathematical
 - \$and, \$or, \$eq, \$lt, \$lte, \$gt, \$gte, \$cmp, \$cond, \$switch, \$in, etc
- Date
 - \$dateFromParts, \$dateToParts, \$dateFromString, \$dateToString, \$dayOfMonth, \$isoWeek, \$minute, \$month, \$year, etc.
- String
 - \$toUpperCase, \$toLowerCase, \$substr, \$strcasecmp, \$concat, \$split, etc.
- Literals
 - \$exp, \$let, \$literal, \$map, \$type, etc

MongoDB Replica Sets



Replica Set – 2 to 50 copies

Self-healing

Data Center Aware

Addresses availability considerations:

- High Availability
- Disaster Recovery
- Maintenance

Workload Isolation: operational & analytics


Serving global audiences with relational databases



MongoDB's native replica sets puts your entire database right next to users



Scaling MongoDB: Automatic Sharding




Multiple sharding policies: hashed, ranged, zoned

Increase or decrease capacity as you go


Automatic balancing for elasticity

Co-locating operational and analytical workloads


Operational
(interactive)




PRIMARY



Secondary




Secondary




BI & Reporting
mongoDB® Charts



Analytics




Secondary
{use = analytics}



Predictive Analytics



Freedom to run anywhere



- Database that runs the same everywhere
- Leverage the benefits of a multi-cloud strategy
- Global coverage
- Avoid lock-in

Convenience: same codebase, same APIs, same tools, wherever you run




Easy: Drivers and Frameworks


Drivers




Frameworks



2 choices as you move to the cloud: Self-Managed or DBaaS




2 choices as you move to the cloud: Self-Managed or DBaaS



1. Provision instances and storage
2. Configure HA
3. Configure security
4. Configure backup/restore
5. Monitoring & alerting
6. Ongoing upgrades & maintenance

2 choices as you move to the cloud: Self-Managed or DBaaS



1. Provision instances and storage
2. Configure HA
3. Configure security
4. Configure backup/restore
5. Monitoring & alerting
6. Ongoing upgrades & maintenance

**Choose instance, hit deploy,
available in a couple of
minutes**

Atlas

unlocks **agility**
and **reduces**
cost



Self-service and
elastic



Global and highly
available



Secure by default



Comprehensive
monitoring



Managed backup



Cloud agnostic

Atlas Demo



Atlas Workshop



Sign Up for Free

<https://www.mongodb.com/cloud/atlas>

The screenshot shows the MongoDB Atlas sign-up process. On the left, there's a large heading "MongoDB Atlas" and a paragraph describing the service. Below that is a green button labeled "Get started free". To the right, a modal window titled "Cloud Provider & Region" is open. It asks to choose a cloud provider and a region. Under "Cloud Provider", "aws" is selected. Under "Region", "N. Virginia (us-east-1)" is selected. The modal also contains sections for "Configure a free tier cluster" and "Configure cross-region replication (M10 and up)". At the bottom of the modal, there's a toggle switch for "Cross-region replication" which is currently off.



EXERCISE 1: MONGODB ATLAS & USERS

AFTER EXERCISE 1, YOU SHOULD BE ABLE TO...

- Deploy a MongoDB Atlas Cluster (free-tier M0 or paid M10+)
(Go to <https://www.mongodb.com/cloud/atlas>)
- Secure the cluster, e.g. user and IP whitelisting
- Understand the basic features of Atlas

EXERCISE 1: LET'S DO IT!

- Deploy M0 instance on your choice of cloud provider
- Give the cluster a name
- Create a user via Atlas (db.createUser())
 - testuser / testuser



EXERCISE 2: MONGODB COMPASS & CRUD

AT THE END OF THIS EXERCISE 2, YOU SHOULD BE ABLE TO...

- Launch *Compass* (tool to explore, visualize, optimize, and modify MongoDB data) and connect to that cluster
- Use Compass to create a database, create a collection, and import a data set into that collection
- Use Compass to create, read, update, and delete (CRUD) documents in a collection

EXERCISE 2: LET'S DO IT!

- Download and Install Compass: <https://bit.ly/2S2jozr>
- Connect to Atlas with Compass
- Create a database and collection
 - Database: movies
 - Collection: movies
- Download data: <https://bit.ly/2MpsvoJ>
- Import data using Compass

CONNECTION STRINGS

Key	Value
Server	<<insert your server details here>>
SRV Record	Yes on Switch
Authentication	Username / Password
Username	<<insert your username that you created>>
Password	<<insert your password that you created>>
Auth DB	admin

CRUD: LET'S CREATE...

```
db.movies.insertOne(  
{  
    "title": "Ghostbusters",  
    "Year": 1984,  
    "Rated": "PG",  
    "Runtime": 105,  
    "Type": "movie",  
    "Genres": ["comedy", "action"] ,  
    "Director": "Ivan Reitman",  
    "Writers": ["dan aykroyd", "Harold ramis"]  
}  
)
```

SCHEMA AND DATA TYPES

(COMPASS)

CRUD: LET'S READ...

```
db.movies.findOne()  
db.movies.find().pretty()
```

CRUD: LET'S UPDATE...

```
db.movies.updateOne(  
  {  
    title: "Ghostbusters"  
  },  
  {  
    $set: {  
      imdb: { id: "tt0087332", rating: 7.8, votes: 312798 }  
    }  
  }  
)
```

CRUD: LET'S DELETE...

```
db.movies.deleteOne(  
  {  
    title: "Ghostbusters"  
  }  
)
```

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	
“Comedy” as the only genre	
“Comedy” or “Drama”	
“Comedy” and “Drama”	
IMDB Rating>8.0 and PG Rating	
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	
“Comedy” or “Drama”	
“Comedy” and “Drama”	
IMDB Rating>8.0 and PG Rating	
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	db.movies.find({genres:["Comedy"]})
“Comedy” or “Drama”	
“Comedy” and “Drama”	
IMDB Rating>8.0 and PG Rating	
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	db.movies.find({genres:["Comedy"]})
“Comedy” or “Drama”	db.movies.find({genres:{\$in:["Comedy", "Drama"]}})
“Comedy” and “Drama”	
IMDB Rating>8.0 and PG Rating	
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	db.movies.find({genres:["Comedy"]})
“Comedy” or “Drama”	db.movies.find({genres:{\$in:["Comedy", "Drama"]}})
“Comedy” and “Drama”	db.movies.find({ genres: { \$all: ["Comedy", "Drama"] } })
IMDB Rating>8.0 and PG Rating	
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	db.movies.find({genres:["Comedy"]})
“Comedy” or “Drama”	db.movies.find({genres:{\$in:["Comedy", "Drama"]}})
“Comedy” and “Drama”	db.movies.find({ genres: { \$all: ["Comedy", "Drama"] } })
IMDB Rating>8.0 and PG Rating	db.movies.find({"imdb.rating" : {\$gt: 8.0}, rated:"PG"})
title starting with “Dr. Strangelove”	

CRUD: FIND ME MOVIES...

QUESTION	ANSWER
From 1987	db.movies.find({year:1987})
“Comedy” as one of their genres	db.movies.find({genres: "Comedy"})
“Comedy” as only genre	db.movies.find({genres:["Comedy"]})
“Comedy” or “Drama”	db.movies.find({genres:{\$in:["Comedy", "Drama"]}})
“Comedy” and “Drama”	db.movies.find({ genres: { \$all: ["Comedy", "Drama"] } })
IMDB Rating>8.0 and PG Rating	db.movies.find({"imdb.rating" : {\$gt: 8.0}, rated:"PG"})
Title starting with “Dr. Strangelove”	db.movies.find({title: {\$regex: '^Dr. Strangelove'}})



EXERCISE 3: INDEXES AND EXPLAIN PLAN

AT THE END OF THIS EXERCISE, YOU SHOULD BE ABLE TO...

- Identify how a query was run using the Explain Plan
- Create an index on a collection
- See how an index affects query plans

INDEXES AND EXPLAIN PLAN

COMPASS

ADDING INDEX

- Rule of thumb (ESR):
 - Equality
 - Sort
 - Range



EXERCISE 4: AGGREGATION FRAMEWORK

AT THE END OF THIS EXERCISE 4 YOU SHOULD BE ABLE TO...


- Have a basic understanding of what the aggregation framework is
- Write some basic queries using the aggregation framework



Aggregations

Advanced data processing pipeline for transformations and analytics

- Multiple stages
- Similar to a unix pipe
 - Build complex pipeline by chaining commands together
- Rich Expressions





Aggregation Features

A feature rich analytical framework

Pipeline Stages

- \$match
- \$group
- \$facet
- \$geoNear
- \$graphLookup
- \$lookup
- \$project
- \$sort
- \$unwind

Operators

- Mathematical
 - \$add, \$abs, \$subtract, \$multiply, \$divide, \$log, \$log10, \$stdDevPop, \$stdDevSam, \$avg, \$sqrt, \$pow, \$sum, \$zip, etc.
- Array
 - \$push, \$reduce, \$reverseArray, \$addToSet, \$arrayElemAt, \$slice, etc.
- Conditionals
 - \$and, \$or, \$eq, \$lt, \$lte, \$gt, \$gte, \$cmp, \$cond, \$switch, \$in, etc
- Date
 - \$dateFromParts, \$dateToParts, \$dateFromString, \$dateToString, \$dayOfMonth, \$isoWeek, \$minute, \$month, \$year, etc.
- String
 - \$toUpperCase, \$toLowerCase, \$substr, \$strcasecmp, \$concat, \$split, etc.
- Laterals
 - \$exp, \$let, \$literal, \$map, \$type, etc

**LET'S PLAY WITH AGGREGATION
PIPELINE BUILDER IN COMPASS**

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	\$match {genres: "Comedy"}

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	<code>\$match {genres: "Comedy"}</code>
How can you use \$unwind to create an individual document for each country?	

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	\$match {genres: "Comedy"}
How can you use \$unwind to create an individual document for each country?	\$unwind {path: "\$countries"}

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	\$match {genres: "Comedy"}
How can you use \$unwind to create an individual document for each country?	\$unwind {path: "\$countries"}
How can you use \$group to count all the comedies grouped by country?	

AGGREGATION WORKSHOP

QUESTION	ANSWER
How can you use \$match to find all comedies?	<code>\$match {genres: "Comedy"}</code>
How can you use \$unwind to create an individual document for each country?	<code>\$unwind {path: "\$countries"}</code>
How can you use \$group to count all the comedies grouped by country?	<code>\$group { _id: "\$countries", count: {\$sum:1} }</code>

Lunch Break



MongoDB Stitch



A bit of background on why...

More businesses driven by Web/Mobile Apps

- Mobile internet usage has surpassed desktop
 - 56% of usage from Mobile/Tablet
- Apps are expected to work across web, mobile, etc.
 - 58% of apps are developed for multiple platforms
- Apps must be highly available, globally distributed
 - 48% of apps are deployed to the cloud

Creating/Integrating services drives development

- More is expected of developers
 - 64% of developers identified as full-stack
 - Leading to more time spent on integration/security
- Microservice architecture gaining traction
 - 69% of teams considering MSA for new projects
- Too much time spent on maintenance/integration
 - 41% of time spent maintaining
 - Only 39% on new projects

It's easier to develop with MongoDB & Atlas, but building applications is still tough because...

- Applications still need somewhere for the app logic to live


This comes will all the same headaches of maintaining databases the traditional way

- Applications are expected to have features like
 - Facebook sign in
 - Push notifications
 - Email confirmations
 - Robust offline features


Low to No Value

Non Differentiating

Focus Your Energy Where You Can Make a Difference



Save weeks of development and thousands of lines of code



Without Stitch	With Stitch
Provision backend server	Not needed
Install backend runtime environment	Not needed
Code user authentication	{Simple JSON config}
Code data access controls	{Simple JSON config}
Code against each external service API	Code against a single SDK/API
Poll the database for changes	Not needed
Code REST API for frontend to use backend	Not needed
Add code to make backend HA	Not needed
Add code to scale backend	Not needed
Code backend application logic	Provide code for Stitch Functions
Code application frontend	Code application frontend using single SDK
Monitor & manage backend infrastructure	Not needed

MongoDB Stitch Serverless Platform – Services



Stitch QueryAnywhere

Brings MongoDB's rich query language safely to the edge

iOS, Android, Web, IoT



Stitch Functions

Integrate microservices + server-side logic + cloud services

Build full apps, or Data as a Service through custom APIs



Stitch Triggers

Real-time notifications let your application functions react in response to database changes, as they happen



Stitch Mobile Sync


Automatically synchronizes data between documents held locally in MongoDB Mobile and your backend database

(coming soon)

Streamlines app development with simple, secure access to data and services from the client with thousands of lines less code to write and no infrastructure to manage – getting your apps to market faster while reducing operational costs.

ToDo App Demo

- Google Authentication
- Twilio, for SMS integration
- Send text messages to add items to my ToDo list.



MongoDB Stitch Workshop



Workshop



Basic Blog

Workshop - Basic Blog

- What will we build?
 - a blog with a comments section backed by Stitch
- Prerequisites
 - MongoDB Atlas Account
 - MongoDB cluster hosted in Atlas
- Estimated Time to Complete: ~15 minutes



Workshop - Basic Blog - Step 1


- Open a text editor
 - create **blog.html**

```
<html>
  <head>
  </head>
  <body>
    <h3>This is a great blog post</h3>
    <div id="content">
      I like to write about technology because I want to get on the
      front page of hacker news.
    </div>
    <hr>
    <div id="comments"></div>
  </body>
</html>
```

Workshop - Basic Blog - Step 2

- Create a Stitch Application

1. Click **Stitch Apps** in the left-hand corner navigation of the Atlas console
2. Click **Create New Application**
3. Give the application a name (e.g. **Blog Tutorial**)
4. Click **Create**
5. Wait for your application to initialize



Workshop - Basic Blog - Step 3

- Turn on Anonymous Authentication
 - From the **Getting Started** page
 - Enable **Anonymous Authentication** under the **Turn On Authentication** heading

Workshop - Basic Blog - Step 4

- Configure **blog.comments** MongoDB Collection
1. Click **Rules** under MongoDB Atlas in the left-hand navigation
 2. Click **Add Collection**
 3. Enter **blog** for the **Database Name**
 4. Enter **comments** for the **Collection Name**
 5. Select **No Template**
 6. Click **Add Collection**

Workshop - Basic Blog - Step 5

- Select the **Permissions** tab of the rules for the collection
 - enable reading and writing to the **comments** collection

1. Click the **Read** and **Write** checkboxes for the **default** role
2. Click **Save**

The screenshot shows the MongoDB Atlas Rules interface. The top navigation bar has tabs for 'Rules' and 'Cluster Info'. Below the navigation, there's a sidebar with a tree view showing a 'blog' collection containing a 'comments' sub-collection. The main area is titled 'blog.comments' and contains tabs for 'Permissions*', 'Schema', and 'Filters'. Under the 'Permissions*' tab, there's a table for the 'default' role. In the 'Fields' section of the table, both the 'Read' and 'Write' checkboxes are checked. There are also sections for 'All Additional Fields' with checked checkboxes. At the bottom right of the main area, there are 'DISCARD CHANGES' and 'SAVE' buttons. A green banner at the bottom of the interface says 'Arrange roles by the order they will be applied' and 'ADVANCED MODE'.

Workshop - Basic Blog - Step 6

- Add a Commenting System to the Blog Post
 - add the following `<script>` tag to the `<head>` section of `blog.html`

```
<script src="https://s3.amazonaws.com/stitch-sdks/js/bundles/4.0.0/stitch.js"></script>
```

Workshop - Basic Blog - Step 7

- Initialize Stitch Clients
 - Paste the following <script> tag in `blog.html` beneath the tag that imports Stitch SDK

```
<script>
  // Initialize the App Client
  const client = stitch.Stitch.initializeAppClient("<your-app-id>");

  // Get a MongoDB Service Client
  const mongodb = client.getServiceClient(
    stitch.RemoteMongoClient.factory,
    "mongodb-atlas"
  );

  // Get a reference to the blog database
  const db = mongodb.db("blog");
</script>
```

Replace <your-app-id> with your **Stitch App ID**

* find your **App ID** on the **Clients** page of the Stitch UI

Workshop - Basic Blog - Step 8

- Query and Display Comments on Page Load
 - Add the following function to the <script> tag

```
function displayComments() {  
  db.collection("comments")  
    .find({}, {limit: 1000})  
    .asArray()  
    .then(docs => {  
      const html = docs.map(doc => `<div>${doc.comment}</div>`);  
      document.getElementById("comments").innerHTML = html;  
    })  
}
```

Workshop - Basic Blog - Step 9

- Log In and Display Comments On Load
 - Setup anonymous authentication in `blog.html`
 - Add the following function to the `<script>` tag:

```
function displayCommentsOnLoad() {  
  client.auth  
    .loginWithCredential(new stitch.AnonymousCredential())  
    .then(displayComments)  
    .catch(console.error);  
}
```

- Update the `<body>` tag in `blog.html` with

```
<body onload="displayCommentsOnLoad()">
```

Workshop - Basic Blog - Step 10

- Add a Form for Submitting Comments
 - Add the `addComment()` function to the `<script>` tag

```
function addComment() {  
  
  const newComment = document.getElementById("new_comment");  
  
  console.log("add comment", client.auth.user.id)  
  
  db.collection("comments")  
    .insertOne({ owner_id : client.auth.user.id, comment: newComment.value })  
    .then(displayComments);  
  
  newComment.value = "";  
}
```

- Add a comment form in the `<body>` of `body.html`

```
<hr>  
  
Add comment:  
  
<input id="new_comment"><input type="submit" onClick="addComment()">
```



Workshop - Basic Blog - Summary

At this point, you should be able to refresh the page and submit a comment on the blog post.

This is a great blog post

I like to write about technology because I want to get on the front page of hacker news.

hello world

,
Lots of snow today!

Add comment:

Workshop



Dashboard

Workshop - Dashboard

- What will we build?
 - a real-time dashboard that displays customer purchase data from the point-of-sale system of a hypothetical pizza restaurant
 - The dashboard will have an automatically updating chart of recent customer purchases and a table of the most popular toppings from the last 100 pizza orders
- Prerequisites
 - ✓ **MongoDB Atlas** Account
 - ✓ MongoDB cluster hosted in Atlas
 - A MongoDB Atlas User **API Key** for your cluster
 - **Node.js** version 6.0.0 or newer
 - A properly installed copy of stitch-cli that has been added to your system PATH
 - A local copy or clone of the stitch-examples GitHub repository
- Estimated Time to Complete: 25 minutes



Workshop - Generate API Key

1. On the upper-right hand corner, click on your user name and select **Account**. Click on **Public API Access**.
2. In the **API Keys** section, click **Generate**.
3. Type a description and click **Generate**.
4. If prompted for two-factor authentication, enter the code and click **Verify**. Then click **Generate** again.
5. **Copy and Record the key immediately**. Atlas displays the full key one time only. You will not be able to view the full key again.
6. Record the key in a secure place. After you record the key, click **Close**.

Account Settings

Profile Personalization Organizations **Public API Access**

API Keys

These are your keys for the Public API. You may have up to ten keys. An API key should be treated like a password, so previously generated keys are only partially shown.



Enabled	Key	Description	Last Used	Created	Actions
✓	*****-***-***-1eda7b1b628e	Workshop	--	01/28/19 - 04:17:57 PM	...

Workshop - Node.js

Node.js is an open-source, cross-platform JavaScript run-time environment that executes JavaScript code outside of a browser.

Download and Install Node.js:

<https://nodejs.org/en/download/>



The screenshot shows the Node.js official website's download page. At the top, there's a navigation bar with links for HOME, ABOUT, DOWNLOADS, DOCS, GET INVOLVED, SECURITY, NEWS, and FOUNDATION. The FOUNDATION link is highlighted with a green background. Below the navigation, the word "Downloads" is centered. A sub-header below it says "Latest LTS Version: 10.15.0 (includes npm 6.4.1)". A descriptive text block encourages users to "Download the Node.js source code or a pre-built installer for your platform, and start developing today." Two main sections are shown: "LTS" (Recommended For Most Users) and "Current" (Latest Features). Under the LTS section, there are links for "Windows Installer" (with a .msi file icon) and "macOS Installer" (with a .pkg file icon). Under the Current section, there is a link for "Source Code" (with a .tar.gz file icon). To the right of these sections, there's a table showing download links for various platforms and architectures. The table has three columns: "32-bit", "64-bit", and "ARM". Rows include: Windows Installer (.msi), Windows Binary (.zip), macOS Installer (.pkg), macOS Binary (.tar.gz), Linux Binaries (x64), Linux Binaries (ARM), and Source Code (with a link to node-v10.15.0.tar.gz).

32-bit	64-bit	
32-bit	64-bit	
64-bit		
64-bit		
64-bit		
ARMv6	ARMv7	ARMv8
node-v10.15.0.tar.gz		

Workshop - Install **stitch-cli**

The easiest way to install **stitch-cli** is with the npm package manager. Ensure that you have Node.js installed, then run the following command in your shell:

```
npm install -g mongodb-stitch-cli
```

Workshop - Get a Local Copy of the Stitch Examples

<https://github.com/mongodb/stitch-examples/>

```
git clone https://github.com/mongodb/stitch-examples.git
```

The screenshot shows the GitHub repository page for 'mongodb / stitch-examples'. The repository has 29 stars and 96 forks. It contains 178 commits, 13 branches, 0 releases, and 23 contributors. The Apache-2.0 license is listed. A pull request button and a 'Clone or download' button are visible. A modal window titled 'Clone with HTTPS' provides the URL <https://github.com/mongodb/stitch-examples>. The repository lists several recent commits:

- danielhelfand and MongoCaleb Add deleteComments Function and Button (#64) ...
- IoTTemperatureTracking fixed various errors in the no_hardware app of the IoT T
- StitchSamplePushNotificationApp DOCSP-2658: Update Android push notifications app t
- basic-blog Add deleteComments Function and Button (#64)
- blog-comments-simple Update JS examples for JS SDK 3.0.0 (#39)

A MongoDB logo is visible in the bottom left corner.

Workshop - Dashboard

dashboard/appdir/dashboard-stitch contains configurations for the following Stitch entities:

- An **email/password authentication provider**.
- A MongoDB service that will link with your Atlas cluster. The service has custom read and write **Rules** defined for the SalesReporting. Receipts collection. The rules remove sensitive fields, such as a customer's credit card number, from read results and require specific fields to be included when inserting new documents.
- A **Stitch Function** named **salesTimeline**.
 - The function returns data from MongoDB for orders made between the provided start and end timestamps.
- A **Stitch Function** named **getPopularToppings**.
 - The function returns a sorted count of the most popular toppings from the last 100 pizza orders stored in MongoDB.
- An **HTTP service** named **PizzaOrderAPI**.
 - The service exposes a webhook named addOrder that adds customer order data to MongoDB when it is sent in the body of a POST request to the webhook. The webhook requires a secret query parameter to validate requests.

Workshop - Dashboard - Step 1

Import a New Stitch Application

1. Navigate to the Pre-Configured Application Directory

```
cd dashboard/appdir/dashboard-stitch
```

1. Authenticate a MongoDB Atlas User

```
stitch-cli login --username=<MongoDB Cloud username> --api-key=<MongoDB  
Atlas API Key>
```

1. Optional: Update MongoDB Atlas Cluster Name.

```
Edit dashboard-stitch/services/mongodb-atlas/config.json
```

1. Import the Application Directory. Run the following command from within the dashboard-stitch directory:

```
stitch-cli import
```



Workshop - Dashboard - Step 2

Create an Email/Password User

1. Open your application in the Stitch UI.
2. Select **Users** from the left-side navigation
3. Click **Add New User**.
4. Enter the following user information then click **Create**:

Username example@email.com

Password mypassword

Workshop - Dashboard - Step 3

Configure the addOrder Webhook to Run as the New User

1. Select **Services** from the left-side navigation
2. Select the **PizzaOrderAPI** service then click the **addOrder** webhook.
3. From the **Settings** tab, set **Run Webhook As** to **User Id** then click the **Select User** button that appears.
4. Choose the email/password user that you just created from the list of users then click **Select User**.
5. Click **Save** to save your changes to the webhook configuration.

Note the values of the **Webhook URL** and **Secret**.

The screenshot shows the MongoDB Stitch Dashboard interface. On the left, there's a sidebar with 'Stitch Apps' and 'DashboardTutorial' selected. Under 'DashboardTutorial', the 'APP ID' is 'dashboardtutorial-ph'. The main area shows the 'addOrder' webhook configuration under the 'Settings' tab. The 'Webhook URL' field contains a redacted URL, and the 'Webhook Name' field contains 'addOrder'. A toggle switch is turned on for 'Respond With Result'. In the 'Run Webhook As' section, 'User Id' is selected. Below it, 'Execute as a specific user' has 'example@email.com' chosen. The 'HTTP Method' dropdown has 'POST' selected. Under 'Request Validation', 'Require Secret As Query Param' is selected. At the bottom, the 'Secret' field shows a redacted value. A green 'Save' button is visible at the top right.

Workshop - Dashboard - Client

The dashboard client code is located in the `stitch-examples/dashboard` directory. There are two files that communicate with Stitch:

`dashboard.js`

This file contains the dashboard client JavaScript code. It queries data from MongoDB using the **`salesTimeline`** and **`getPopularToppings`** Functions that we imported in the previous section.

`data_generator.js`

This script simulates customer orders being entered into a point-of-sale system at a pizza shop. The order data is uploaded to MongoDB Atlas by including it as the body of a POST request sent to the **`addOrder`** webhook.

Workshop - Dashboard - Step 4

Configure the Data Generation Script

1. In `stitch-examples/dashboard/data_generator.js`, find and replace `<YOUR WEBHOOK>` with the Webhook URL you got from the addOrder configuration page. Make sure to append the secret as in the following prototype:

```
<webhook url>?secret=yummypizza
```

1. Navigate to the `stitch-examples/dashboard` directory and run the following command to install all Node.js dependencies:

```
npm install
```

Workshop - Dashboard - Step 5

Connect the Dashboard Client to your Stitch Applications

1. In `stitch-examples/dashboard/dashboard.js`, find and replace `<YOUR APP ID>` with your Stitch application's App ID.

Run the Data Generation Script `data_generator.js`

1. Run the following command in the dashboard directory to begin uploading simulated order data to your MongoDB Atlas cluster:

```
node data_generator.js
```

Workshop - Dashboard - Step 6


Launch the Dashboard

From a new shell window, run the following command in the dashboard directory to begin serving the dashboard locally:

```
npm start
```

Log In to the Dashboard

Navigate to **localhost:8080** in your browser



The image shows a screenshot of a web browser displaying a login page for a "Sales Dashboard". The page has a light gray background with a white central form area. At the top center, the text "Sales Dashboard" is displayed in a large, bold, black font. Below this, centered text reads "Log in to view the dashboard". The form consists of two input fields: "Email:" followed by an input box containing "example@email.com", and "Password:" followed by an input box containing "mypassword". A blue rectangular button labeled "Log In" is positioned below the password field. The overall design is clean and modern.

Workshop - Dashboard - Summary

Sales Dashboard

Logged in as: example@email.com

Popular Toppings

Topping	# Recent Orders
Pineapple	14
Sausage	12
Bacon	12
Spinach	11
Mushrooms	11
Pepperoni	10
Extra cheese	9
Green peppers	9
Onions	7
Black olives	5

Recent Pizza Sales

A line chart titled "Recent Pizza Sales" showing the total order amount over time. The Y-axis is labeled "Order Total (\$)" and ranges from 0 to 30. The X-axis is labeled "Time of Order Submission" and shows times from 12:44 to 14:45. The data is represented by a teal line with a light teal shaded area underneath, showing a fluctuating trend with peaks around 22-24 and troughs around 16-18.

Time of Order Submission	Order Total (\$)
12:44	17
12:48	18
12:52	16
12:56	22
13:00	17
13:04	21
13:08	19
13:12	23
13:16	20
13:20	22
13:24	17
13:28	21
13:32	19
13:36	23
13:40	20
13:44	22
13:48	17
13:52	21
13:56	19
14:00	23
14:04	20
14:08	22
14:12	17
14:16	21
14:20	19
14:24	23
14:28	20
14:32	22
14:36	17
14:40	21
14:44	19

MongoDB Stitch Tutorials

The tutorials in this section walk through creating real applications with MongoDB Stitch. Follow along to get practical experience working with Stitch, or use one of the completed applications as a starting point for your own idea.

Today:


[Link to tutorial](#)




Basic Blog




Dashboard




Temperature Tracker (IoT Integration)



ToDo App - Web



ToDo App - Android



ToDo App - iOS

Tomorrow:


[Stitch](#) > Tutorials

MongoDB Stitch Tutorials

The tutorials in this section walk through creating real applications with MongoDB Stitch. Follow along to get practical experience working with Stitch, or use one of the completed applications as a starting point for your own idea.



Basic Blog




Dashboard




Temperature Tracker (IoT Integration)



ToDo App - Web



ToDo App - Android




ToDo App - iOS

MongoDB Charts



MongoDB Charts: Create, Visualize, Share

```
{
  "id": {"$oid": "5afbc3dc09c8d2dd5852cf2"},
  "saleDate": {"$date": "2017-11-08T19:06:53.449Z"},
  "items": [
    {
      "name": "envelopes",
      "tags": ["stationery", "office", "general"],
      "price": {"$numberDecimal": "9.83"},  
...  
{"name": "pens",
      "tags": ["office", "writing", "school", "stationery"],
      "price": {"$numberDecimal": "73.62"},  
...  
{"name": "laptop",
      "tags": ["office", "school", "electronics"],
      "price": {"$numberDecimal": "595.72"},  
...  
{"name": "notepad",
      "tags": ["office", "writing", "school"],
      "price": {"$numberDecimal": "34.65"},  
...  
]
  },
  "storeLocation": "Seattle",
  "customer": {
    "gender": "M",
    "age": 45,
    "email": "uga@we.soo",
    "satisfaction": 4
  },
  "couponUsed": false,
  "purchaseMethod": "Online"
}
```



Work with complex data

Connect to data sources securely.
Filter. Sample. Visualize.

Share dashboards and
collaborate

Charts Workshop - Visualizing Movie Details

- What will we build?
 - MongoDB Charts dashboard with a data source containing information gleaned from IMDb and Rotten Tomatoes.
- Prerequisites
 - MongoDB Atlas Account
 - MongoDB cluster hosted in Atlas
- Estimated Time to Complete: ~10 minutes

Charts Workshop - Step 1

Add the Data Collection as a Data Source

1. In MongoDB Atlas, click on **Charts** on the left side
2. Click the **Data Sources** tab
3. Click **New Data Source**
4. Select your Atlas Deployment in your project
5. Click **Connect**
6. Select the `movies.movies` collection
7. Click **Set Permissions**, leave the permissions as the default
8. Click **Publish Data Source**

Charts Workshop - Step 2

Create a New Dashboard


1. Click the **Dashboards** tab
2. Click the **New Dashboard** button
3. Enter the Title: Movie Details
4. Click **Create**

Charts Workshop - Step 3

Create Chart Showing Directors with the Most Awards

1. Click **Add Chart**
2. In the Data Source dropdown, select movies.movies
3. Select Chart Type: **Column / Stacked**
4. X Axis: director
 - Sort By: Aggregated Value, Descending
 - Limit: 10
5. Y Axis: awards.wins
 - Aggregate: sum
6. Series: genres
 - Array Reduction: Unwind Array
7. Title: Directors with Most Awards, Split by Genre
8. Click **Save and Close**

Charts Workshop - Summary



Thank You



APPENDIX

